


Machine Learning for Neuroimaging and Neuroscience

Lesson 6

Brain Computer Interface with Deep Generative Models

Alessandro Crimi

a.crimi@sano.science

 @Dr_alex_crimi

<https://bam.sano.science/>

Outline


Deep Convolutional Generative Adversarial
Network (DCGAN) BigBiGAN

Tricks for more realistic image construction
using GANs Image reconstruction from
brain signals

Evaluation of brain readers


Towards Brain Computer Interface (BCI)

Deconvolution



Pytorch: `torch.nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, padding)`
h: `torch.nn.ConvTranspose2d(in_channels=1, out_channels=1, kernel_size=3, stride=0, padding=2)`


Deep Convolutional GAN (DCGAN)



Generative Adversarial Networks



Ian Goodfellow,
2014




Generator **VS** Discriminator

Min-max Game: minimize the possible loss for a worst scenario

Image credit: <https://sthalles.github.io/intro-to-gans/>


Discriminator



```
# Create the dataset
dataset = dset.ImageFolder(root=dataroot,
                           transform=transforms.Compose([
                               transforms.Resize(image_size),
                               transforms.CenterCrop(image_size),
                               transforms.ToTensor(),
                               transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
                           ]))
```

Normalize image pixel value from [0,1] to [-1,1]


Discriminator – pytorch imple on



```
class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 8 x 8
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input)
```


Generator




```
# Generate batch of latent vectors  
noise = torch.randn(b_size, nz, 1, 1, device=device)
```

Random numbers from
normal distribution with mean 0
and variance 1

Tan
h Normalize to
[-1,1]
Con
sistent with
image pixel
input range



Generator - pytorch implementation




```
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 8 x 8
            nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. (ngf*2) x 16 x 16
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # state size. (ngf) x 32 x 32
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            nn.Tanh()
            # state size. (nc) x 64 x 64
        )

    def forward(self, input):
        return self.main(input)
```

Training GAN – Part 1


Training part 1: train

Generator  Discriminator
or Min-max Game



Training GAN – Part 1

Training part 1: train discriminator



Generator **VS** Discriminator
or Min-max Game


```
## Train with all-fake batch
# Generate batch of latent vectors
noise = torch.randn(b_size, nz, 1, 1, device=device)
# Generate fake image batch with G
fake = netG(noise)
label.fill_(fake_label)
# Classify all fake batch with D
output = netD(fake.detach()).view(-1)
# Calculate D's loss on the all-fake batch
errD_fake = criterion(output, label)
# Calculate the gradients for this batch
errD_fake.backward()
D_G_z1 = output.mean().item()
# Add the gradients from the all-real and all-fake batches
errD = errD_real + errD_fake
# Update D
optimizerD.step()
```

```
## Train with all-real batch
netD.zero_grad()
# Format batch
real_cpu = data[0].to(device)
b_size = real_cpu.size(0)
label = torch.full((b_size,), real_label, device=device)
# Forward pass real batch through D
output = netD(real_cpu).view(-1)
# Calculate loss on all-real batch
errD_real = criterion(output, label)
# Calculate gradients for D in backward pass
errD_real.backward()
D_x = output.mean().item()
```

Training GAN – Part 2

Generator **VS** Discriminator
or Min-max Game


Training part 2: train generator



$G(z)$: batch_size, 64, 64, 3

Ground truth label: 1 (real)

****In discriminator training, the ground truth label: 0 (fake)**



Gradients with respect to fake images but labelled as “real”

Training GAN – Part 2

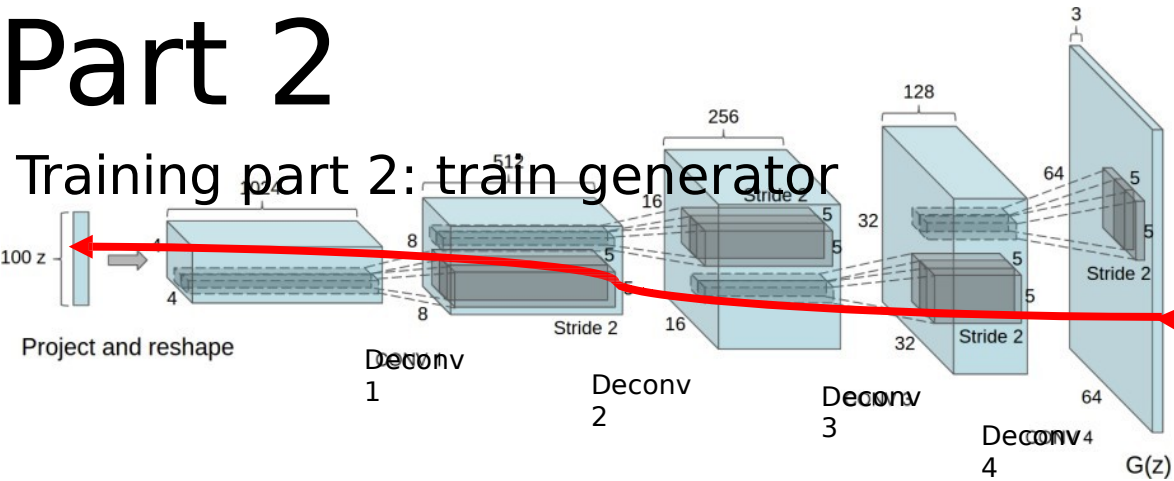
Generator **VS** Discriminator
or Min-max Game

Training part 2: train generator


$G(z)$: batch_size, 64, 64, 3

Ground truth label: 1 (real)

**In discriminator training, the ground truth
Label: 0 (fake)



Gradients with respect to fake images but labelled as "real"



```
#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG.step()
```

GAN Zoo

StyleGAN

cGAN

ProgressiveGAN

CycleGAN

InforGAN

LapGAN

AC-GAN

BigBiGAN

BiGAN

EBGAN

StackGAN


BEGAN

WGAN

BigGAN





Problems with Deep Convolutional GAN (DCGAN)

- Generated images are very small, 64x64, 128x128
- One generator only corresponds with one class of images (no control over random vector z)
- Generated image quality is bad
- Training GAN is brittle:
- Non-convergence: Model parameter oscillate and never converge
- Model collapse: Produce limited number of samples
- Diminished gradients: discriminator is too perfect and generator always fails
- Highly sensitive to hyperparameters



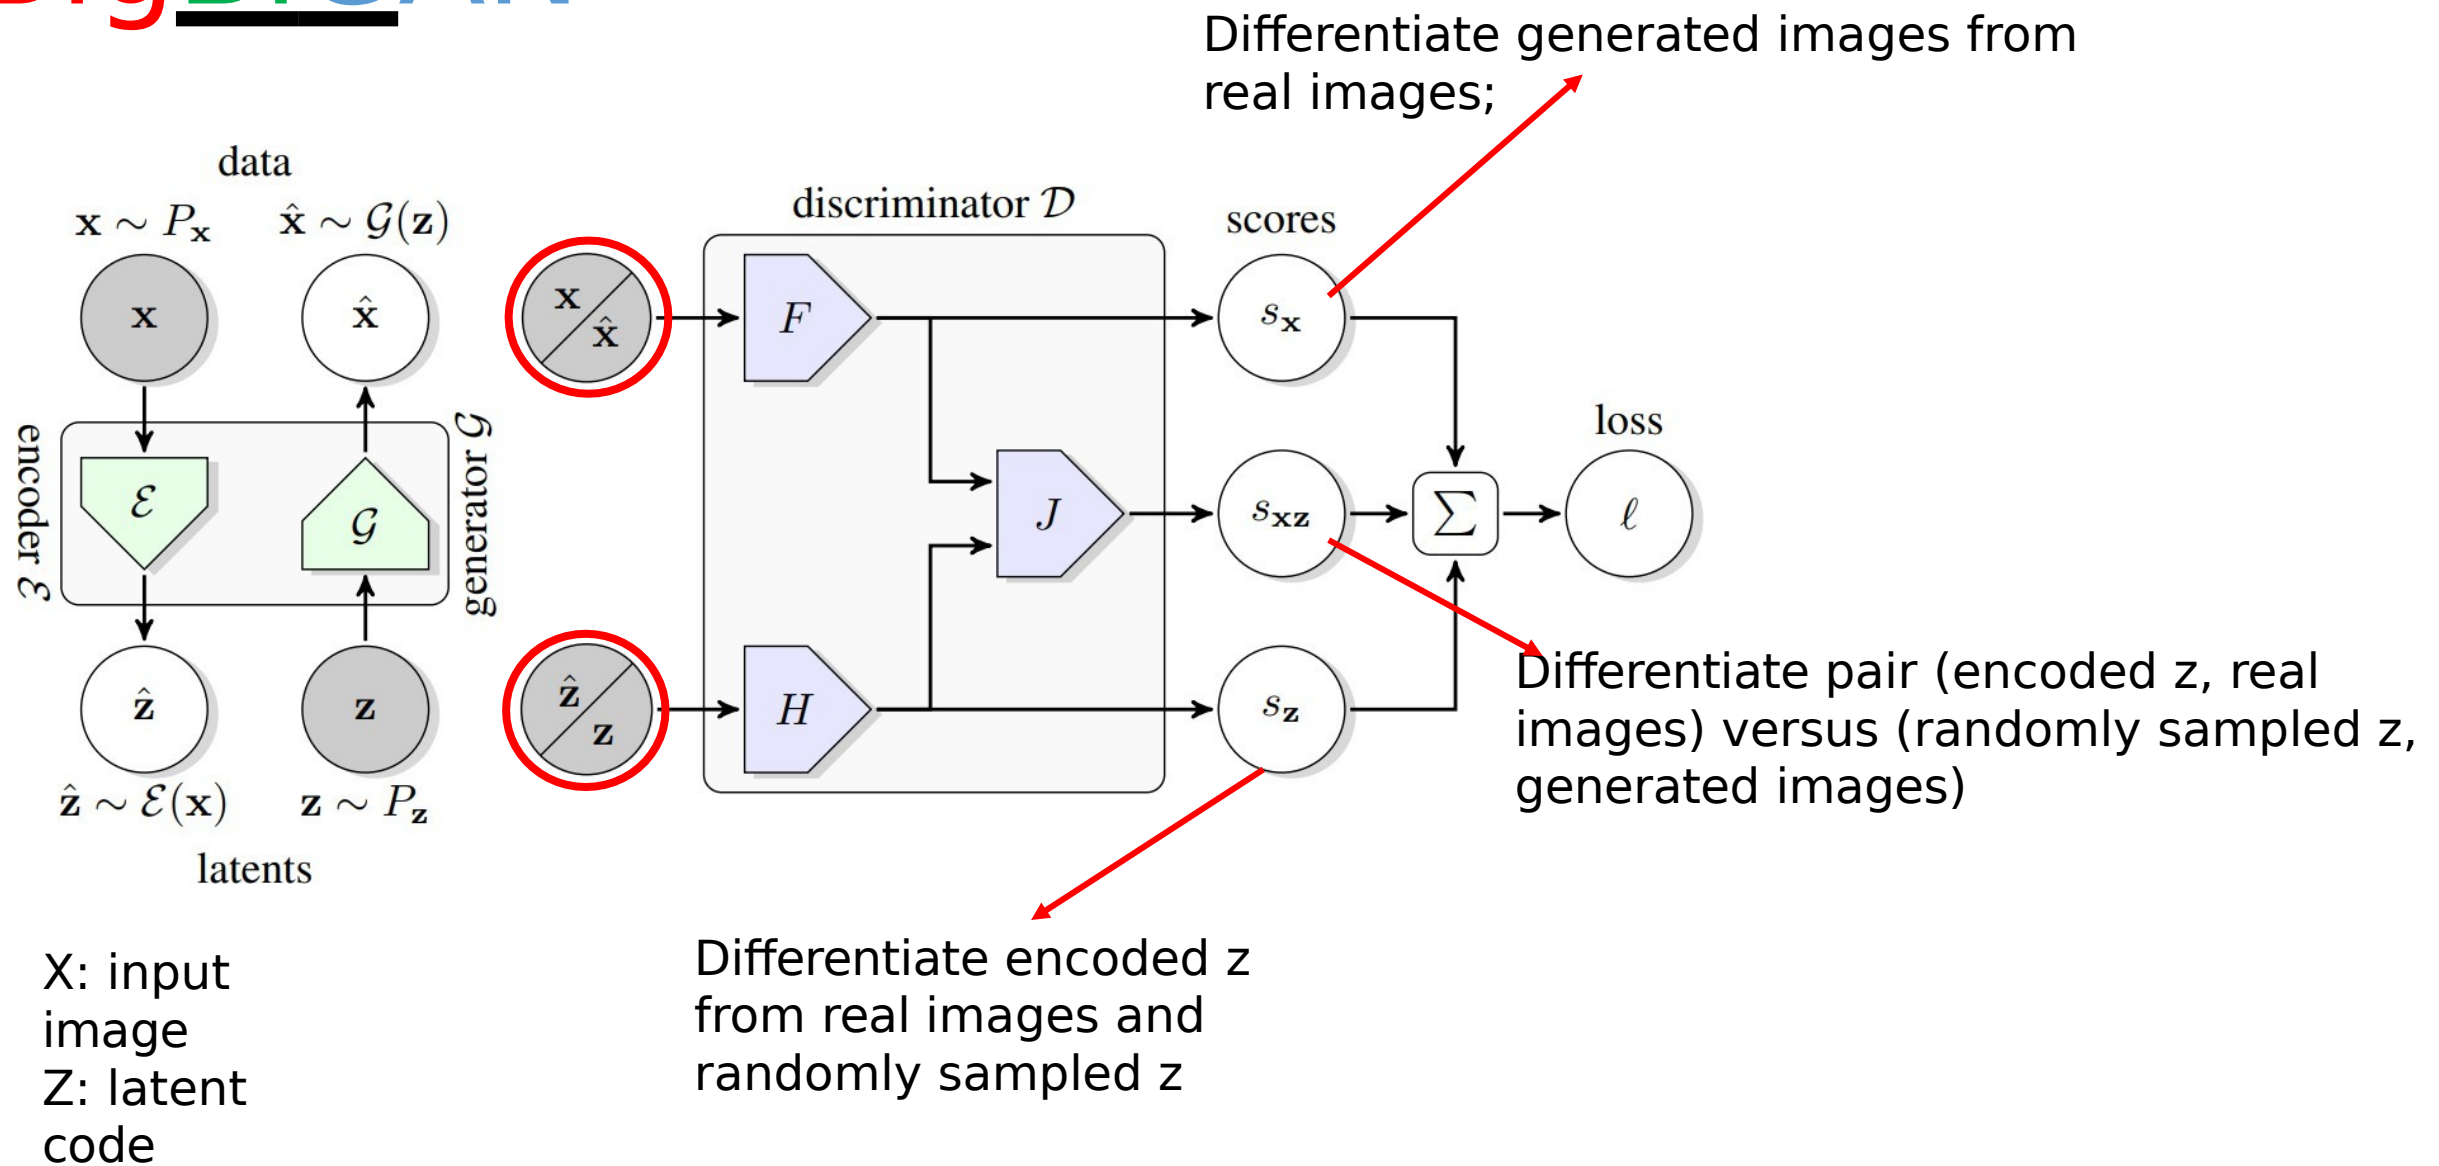
BigBiGAN

Larger batch size, more network parameters, network architecture designs
(skip z-connections, self-attention)





Self-attention Module in GAN

BigBiGAN




Tricks for More Realistic Image Reconstruction


- Update discriminators more often than generators during training



- Moving average of model weights (Progressive GAN)




After 1st epoch,
Generator
Parameter W_{G1}



After 2nd epoch,
Generator
Parameter W_{G2}

...




After Tth epoch,
Generator
Parameter W_{GT}

$$W_{Gfinal} = \text{average}(W_{G1}, W_{G2}, \dots, W_{GT})$$

Tricks for More Realistic Image Reconstruction

- Truncate z resampling at test stage




10
0


Z

```
# Generate batch of latent vectors  
noise = torch.randn(b_size, nz, 1, 1, device=device)
```


- Orthogonal weight initialization
- Orthogonal weight regularization



Brain Reading



Brain Codes from
ECoG, fMRI, EEG,
MEG, etc




All Animal
species

z

10

0



Five senses: Smell,
Images
Touch, Sound, Taste
High-level
functionality:
Emotions,
Memories,
Language

Image Reconstruction Methods from Brain Signals

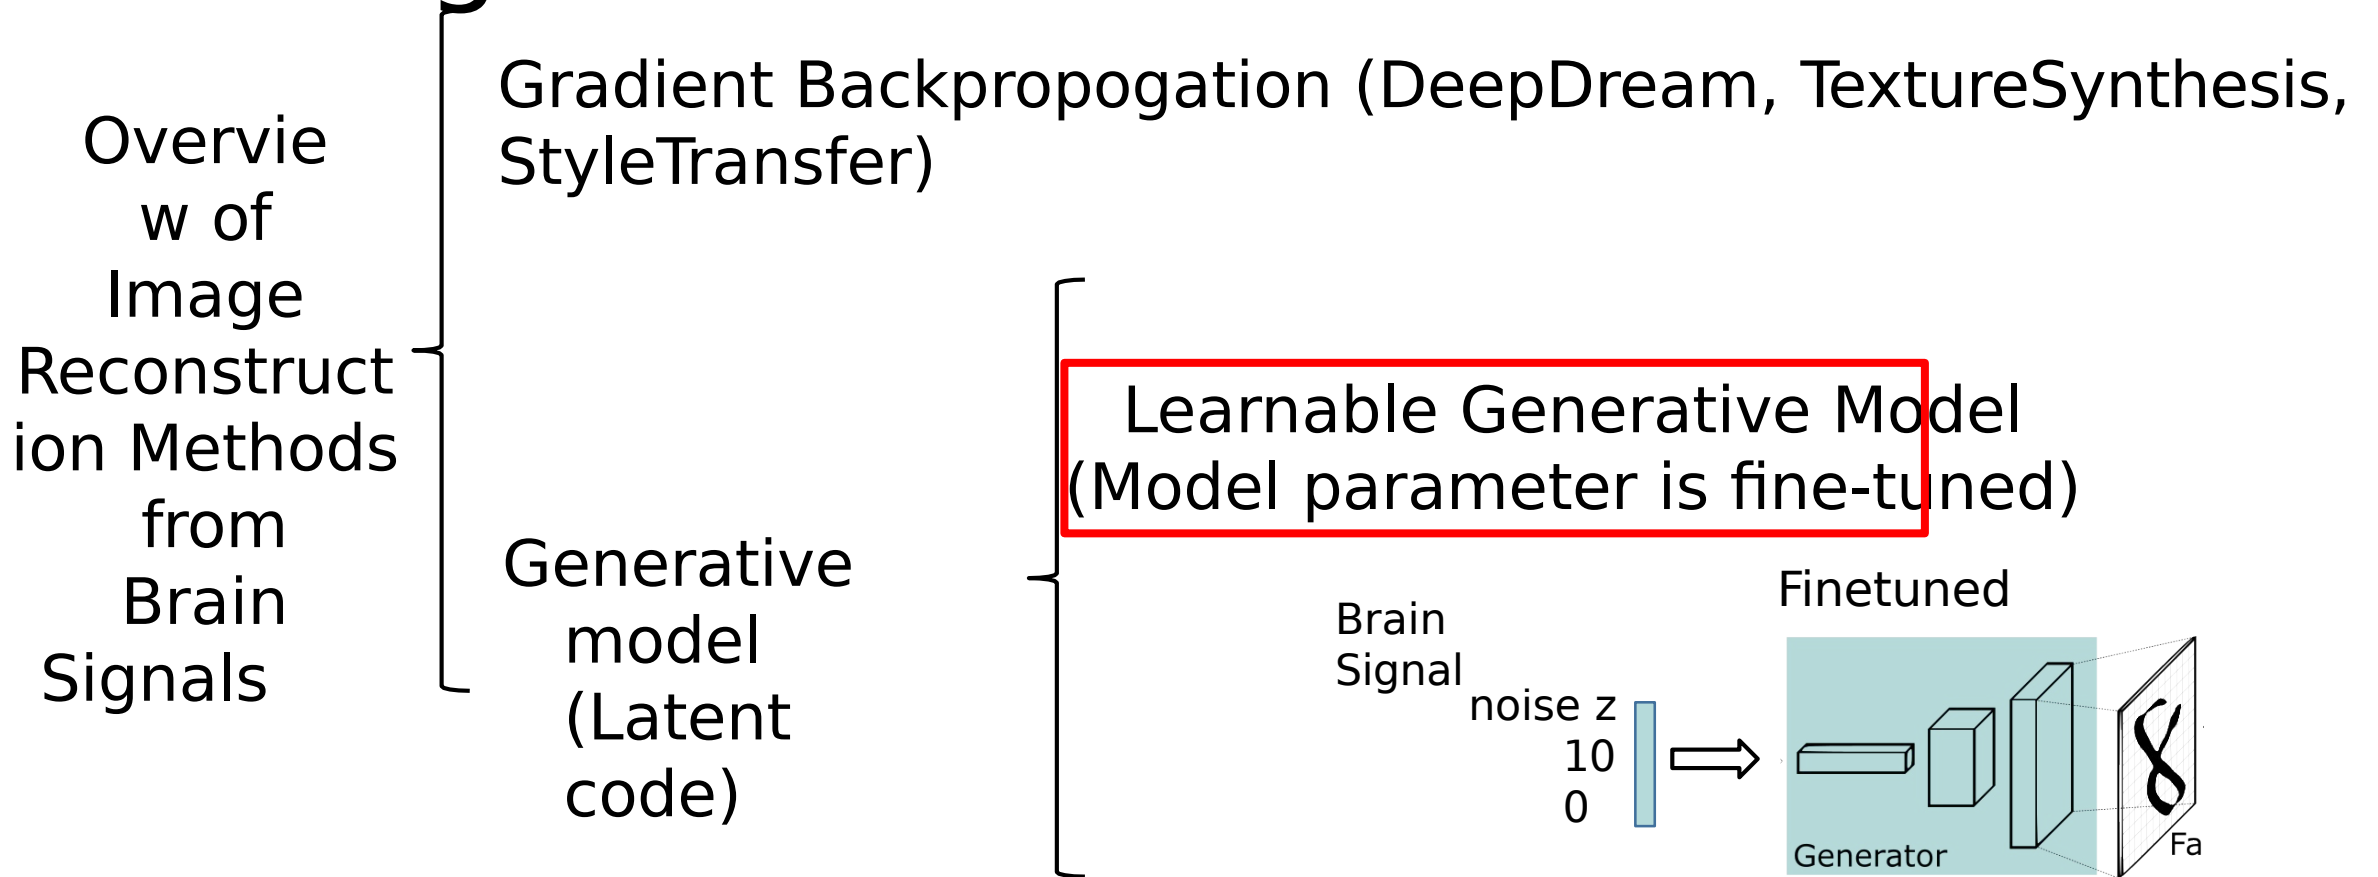
Overview
w of
Image
Reconstruct
ion Methods
from
Brain
Signals

Gradient
Backpropagation

DeepDream

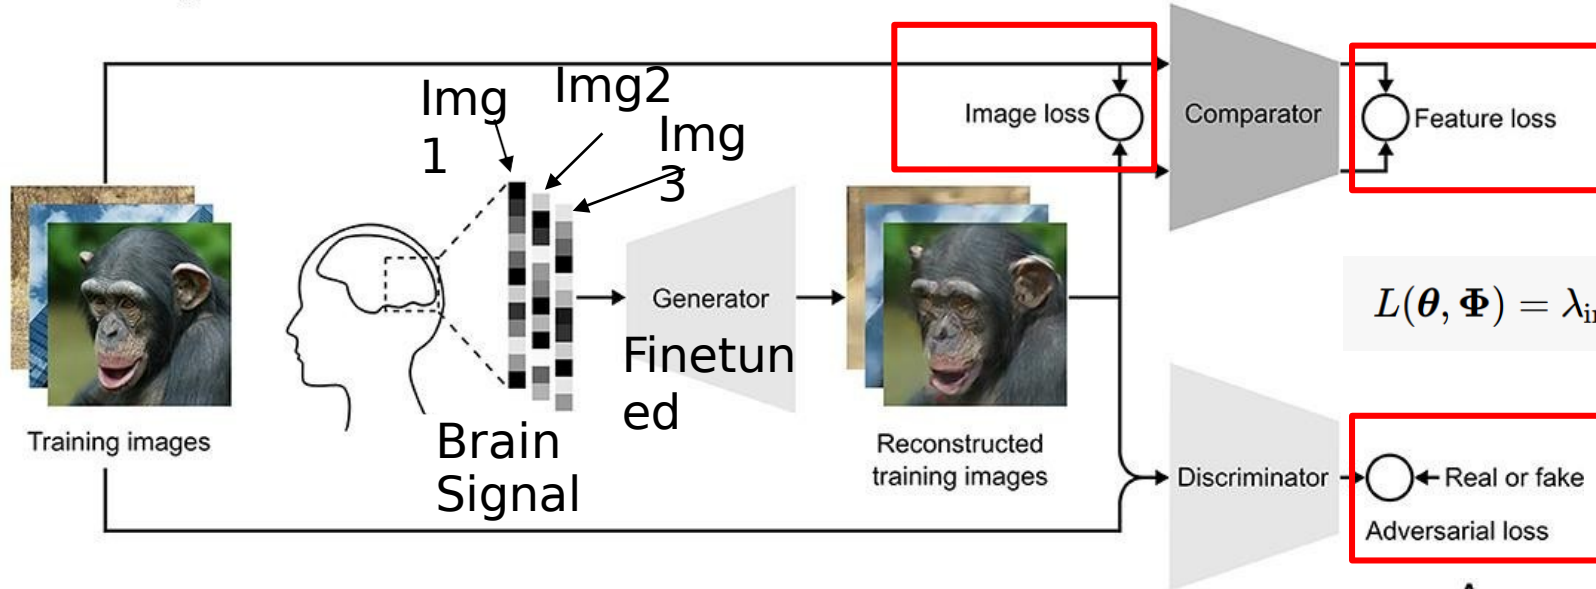
TextureSynthesis,
StyleTransfer

Image Reconstruction Methods from Brain Signals



End-to-End Learnable Generative Model


A Model training



fMRI

$$L(\theta, \Phi) = \lambda_{\text{img}} L_{\text{img}}(\theta) + \lambda_{\text{feat}} L_{\text{feat}}(\theta) + \lambda_{\text{adv}} L_{\text{adv}}(\theta, \Phi)$$

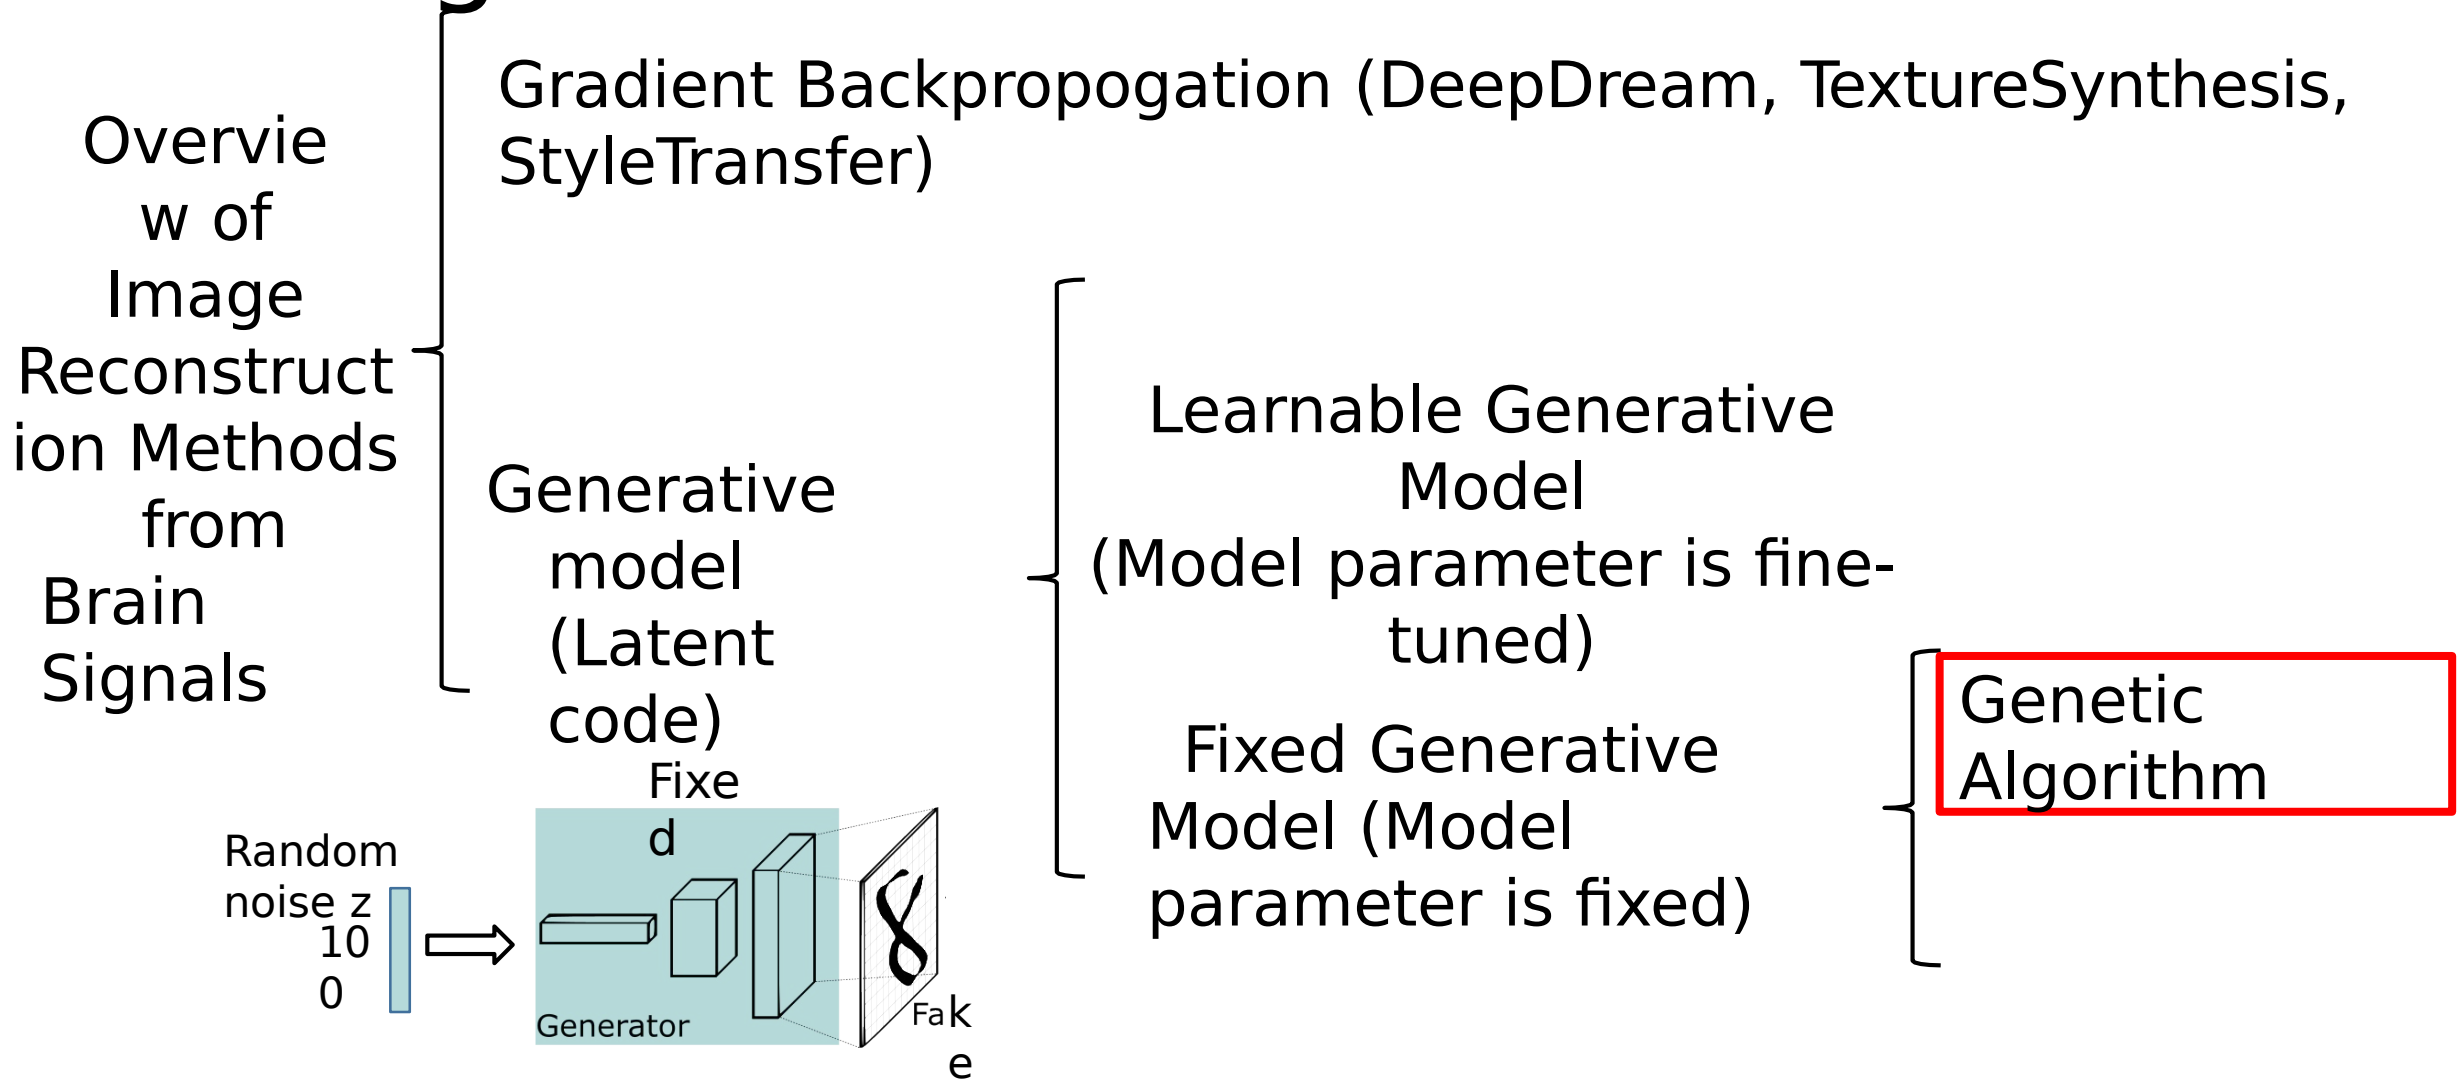
B Model test



A



Image Reconstruction Methods from Brain Signals



Evolving Latent Code using Genetic Alg.

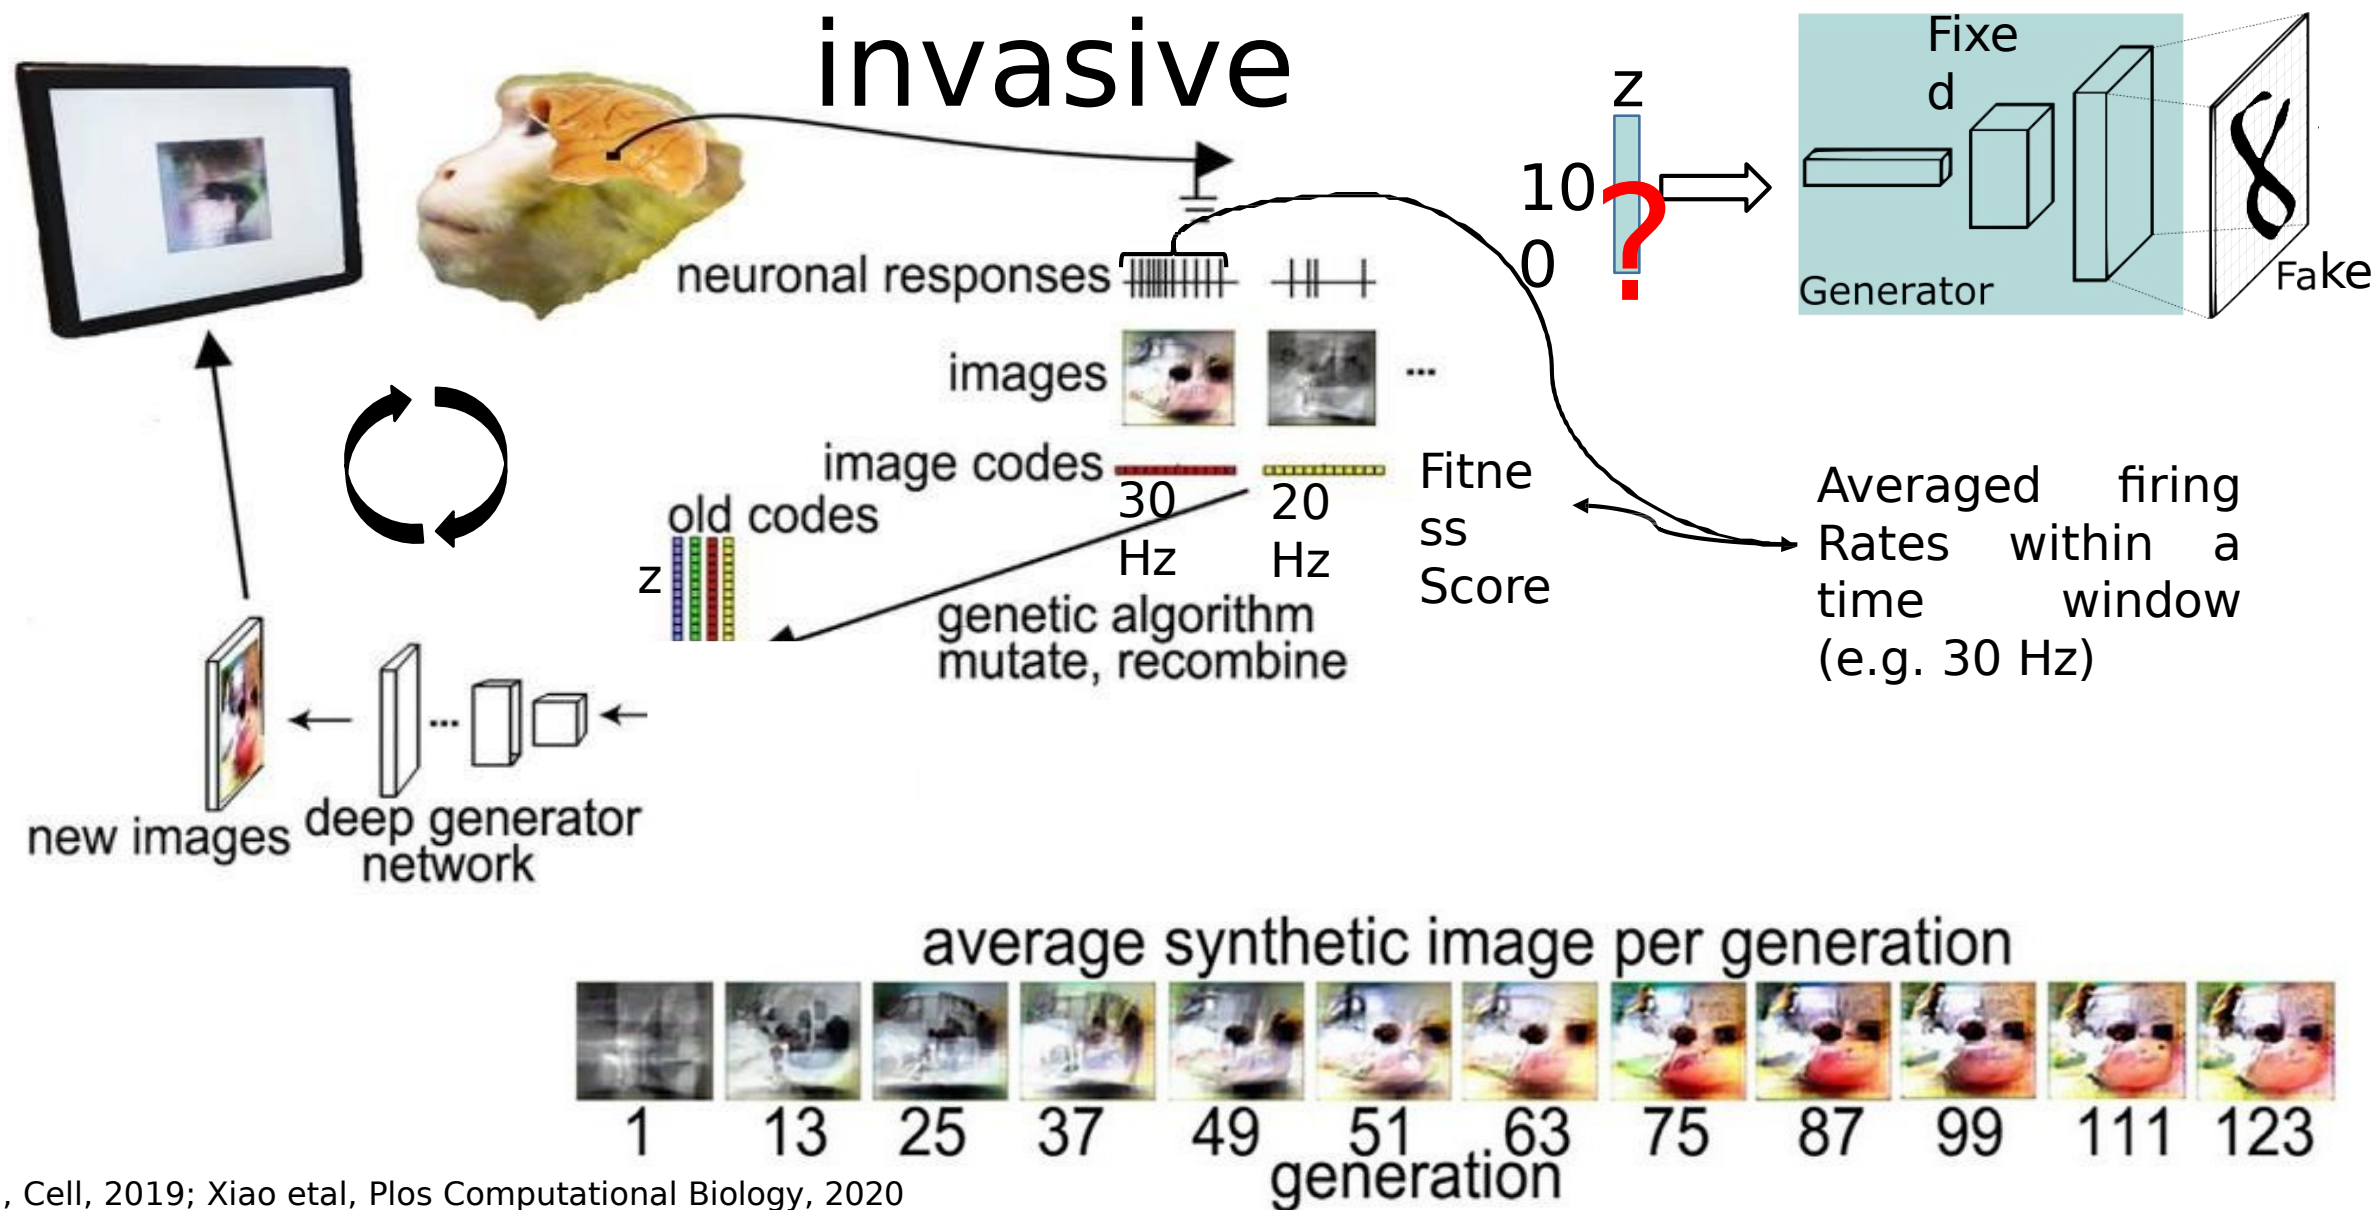




Image Reconstruction Methods from Brain Signals

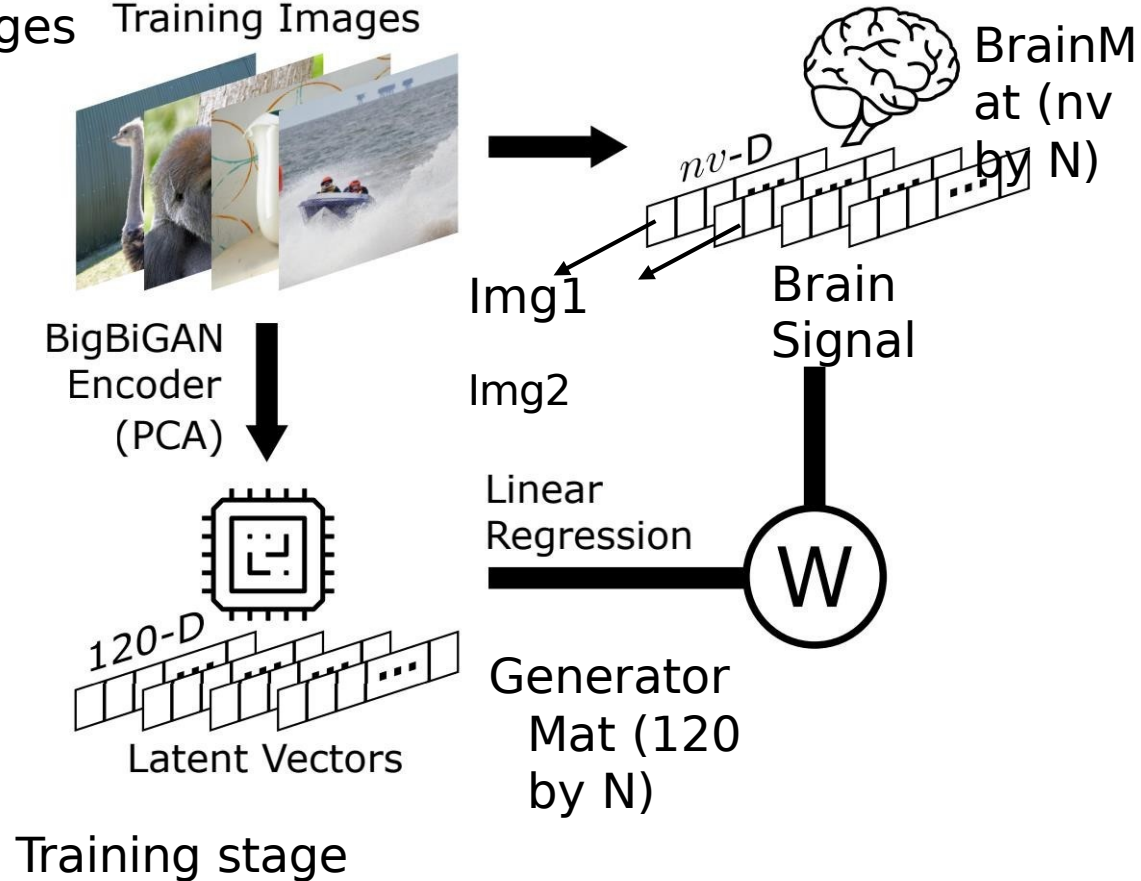


BigBiGAN



Mapping Latent Code using Linear Regression

N total training
images Training Images

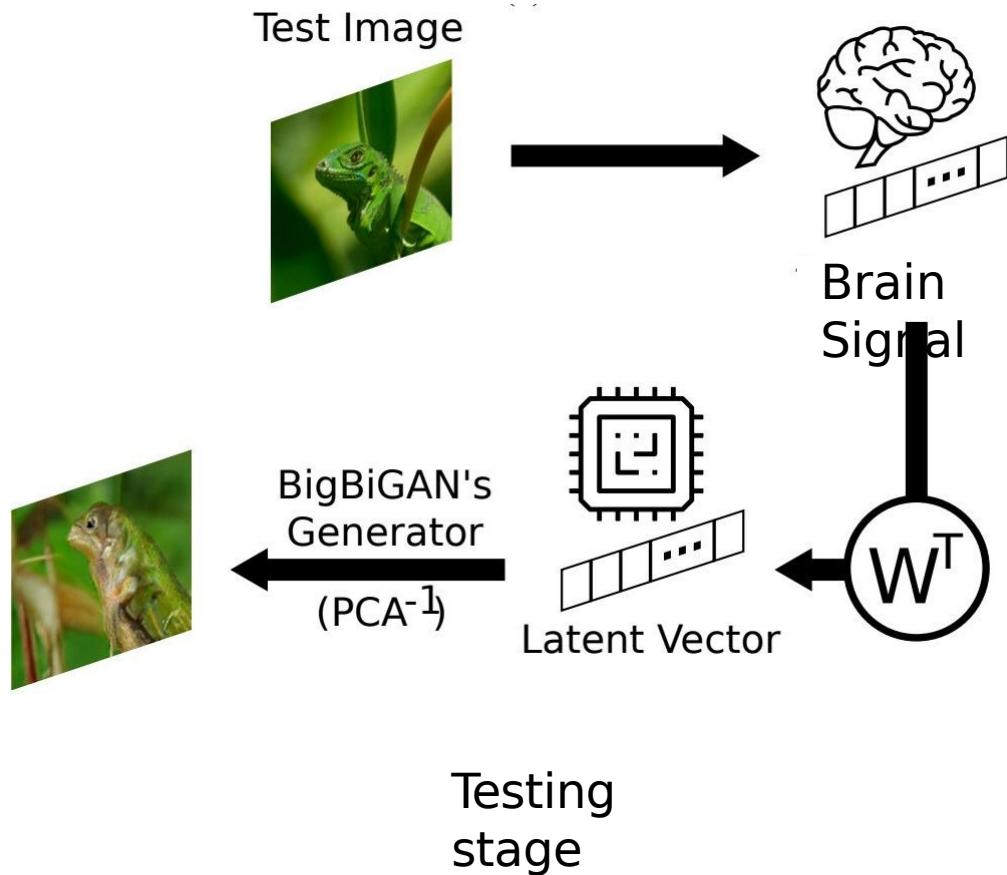



































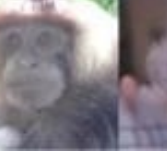
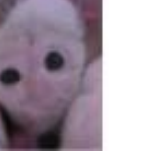
$$\begin{matrix} \mathbf{W} \\ (nv \text{ by } 120) \end{matrix} \text{ Generator Mat } (120 \text{ by } N) = \text{BrainM at } (nv \text{ by } N)$$

Practice:

1. W might be invertible: pseudo-inverse
2. BrainMat size is too large: dimension reduction using PCA to pre-process the data
3. Latent code normalization


Mapping using Linear Regression



#	Input Image	BigBiGAN Recon.	fMRI Reconstruction				
			sub-01	sub-02	sub-03	sub-04	sub-05
1							
2							
3							
4							
5							

Mozafari et al, Reconstructing Natural Scenes from fMRI Patterns using BigBiGAN, 2020

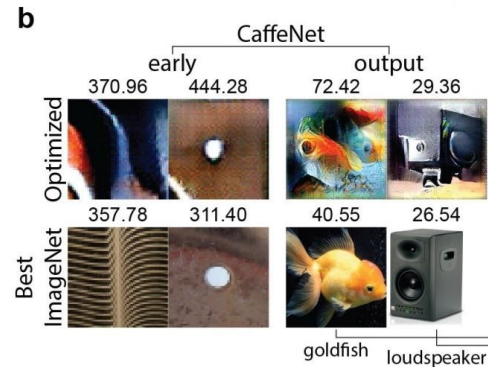
Which brain reader is better?



Bashivan et al,
Science, 2019




Shen et al, Frontiers in Computational
Neuroscience, 2019



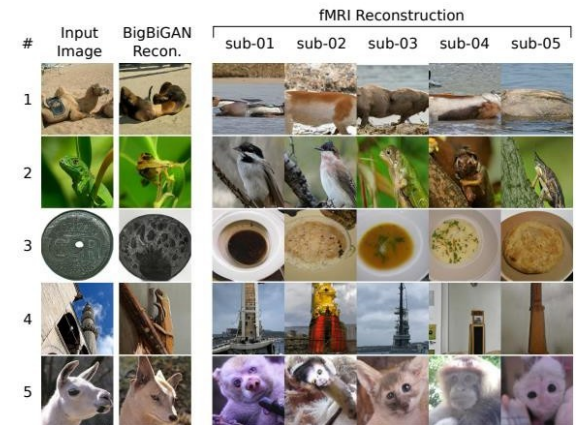
Ponce et al, Cell, 2019;
Xiao et al, Plos Computational Biology,
2020



Shen et al, Plos Computational
Biology, 2019




Cadena et al, Plos Computational
Biology, 2019




Mozafari et al, arxiv,
2020

Do not be subjective. Use quantative metrics to evaluate image reconstruction quality

Towards Brain Computer Interface




Brain Machine Interface (BCI) for the visually impaired



Preliminaries on Genetic Algorithm

Natural selection process where the fittest individuals are selected for reproduction in order to produce offspring of the next generation. -> the fittest latent code z which drives neurons to fire as much as possible




Population, Chromosomes and Genes

		Fitness score (Average Firing Rates)
A1	0 0 0 0 0 0 0	2
A2	1 1 1 1 1 1 1	0
A3	1 0 1 0 1 1 1	4
A4	1 1 0 1 1 1 0	5

Selection
on
A1 0 0 0 0 0 0 0
A2 0 0 0 0 0 0 0

Fittest
parents
A1
2



Before Mutation

A5 1 1 1 0 0 0 0

After Mutation

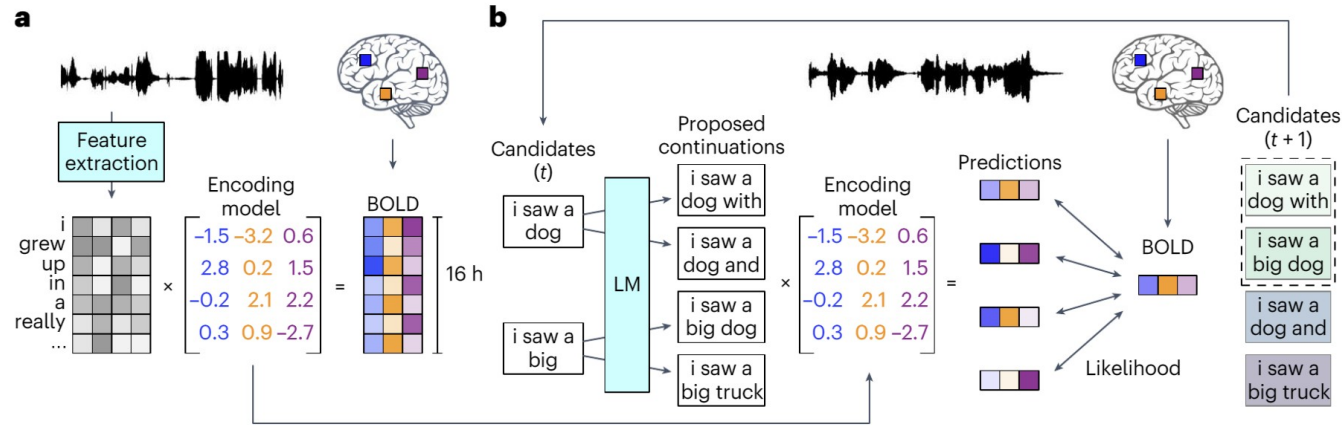
A5 1 1 0 1 1 0 0

Mutation: Before and After

```

START
Generate the initial population
Compute fitness
REPEAT
    Selection
    Crossover
    Mutation
    Compute fitness
UNTIL population has converged
STOP
    
```


“SEMANTIC RECONSTRUCTION OF CONTINUOUS LANGUAGE FROM NON-INVASIVE BRAIN RECORDINGS” TANG ET AL. NATURE NEUROSCIENCE 2023



c

Actual stimulus	Decoded stimulus
<i>i got up from the air mattress and pressed my face against the glass of the bedroom window expecting to see eyes staring back at me but instead finding only darkness</i>	i just continued to walk up to the window and open the glass i stood on my toes and peered out i didn't see anything and looked up again i saw nothing
<i>i didn't know whether to scream cry or run away instead i said leave me alone i don't need your help adam disappeared and i cleaned up alone crying</i>	started to scream and cry and then she just said i told you to leave me alone you can't hurt me anymore i'm sorry and then he stormed off i thought he had left i started to cry
<i>that night i went upstairs to what had been our bedroom and not knowing what else to do i turned out the lights and lay down on the floor</i>	we got back to my dorm room i had no idea where my bed was i just assumed i would sleep on it but instead i lay down on the floor
<i>i don't have my driver's license yet and i just jumped out right when i needed to and she says well why don't you come back to my house and i'll give you a ride i say ok</i>	she is not ready she has not even started to learn to drive yet i had to push her out of the car i said we will take her home now and she agreed

Exact

Gist

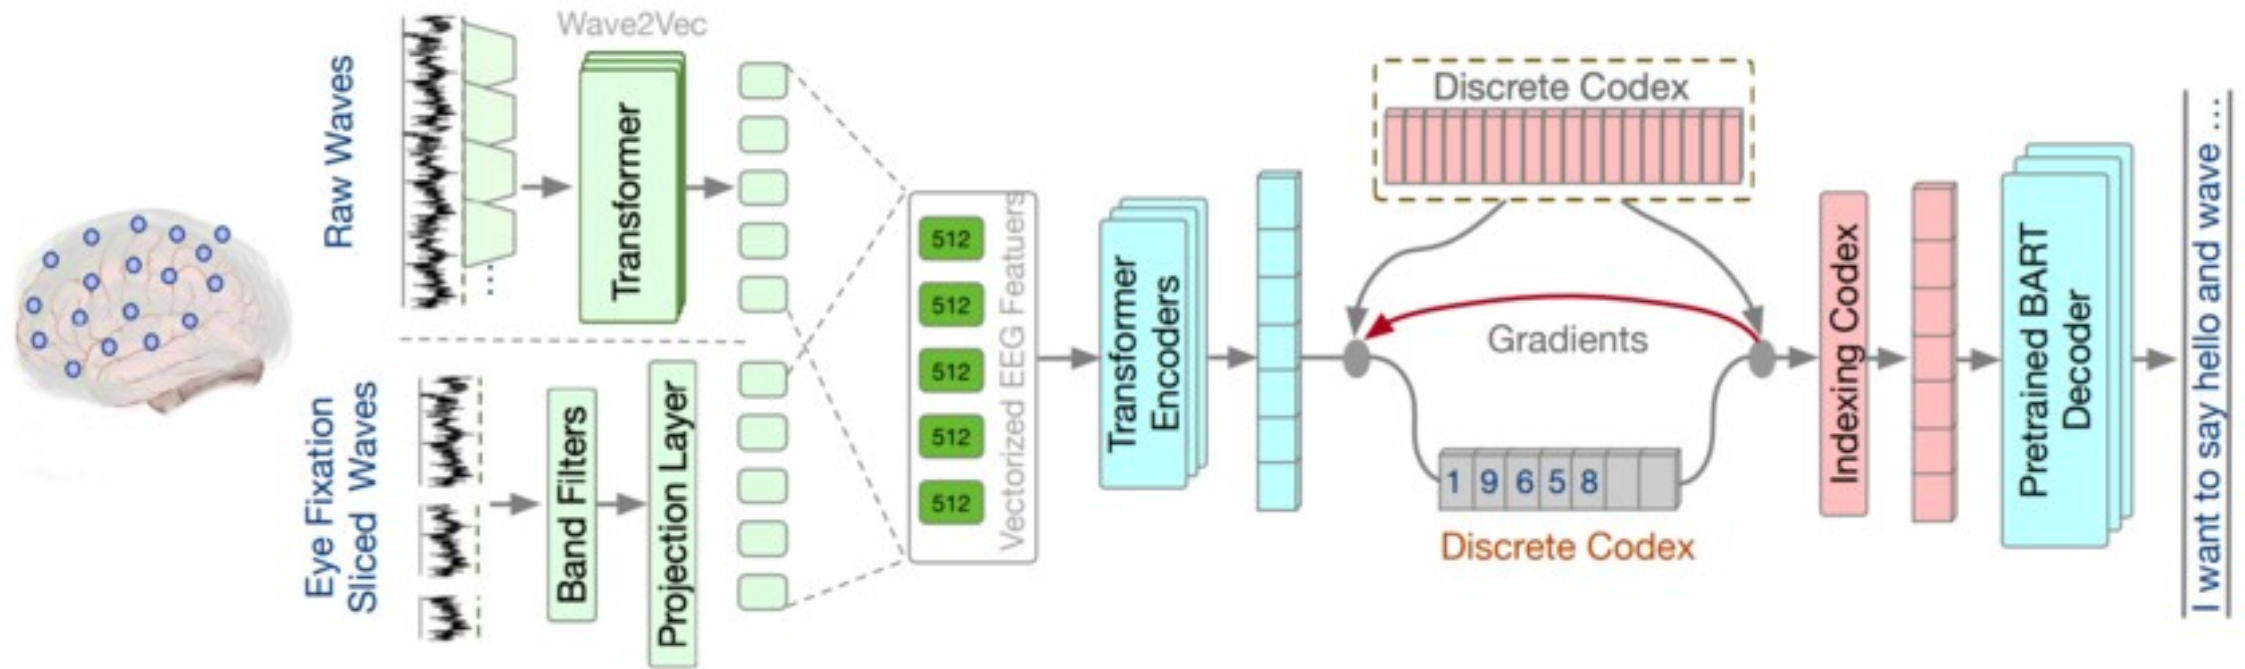
Error

c

Actual stimulus	Decoded	
		she was very weak i held her neck to get her breathing under control
		i see a girl that looks just like me get hit on her back and then she is knocked off

"DeWave: Discrete #EEG Waves Encoding for #Brain Dynamics to Text Translation"

<https://arxiv.org/abs/2309.14030>





Questions?